

D590_Final_Project_Part_2

Brian Kenney, Prabakar Mohan, Brendan Kelly

2023-03-27

Load Library Packages

AAPL

Set up File name

```
file = "Data/AAPL_2023-03-26.csv"
```

Read the .CSV and select Date and Closing Price

```
Stock_information <- read.csv(file)
selected_stock_information <- Stock_information %>% select(c('Date', 'Close')) %>% mutate(row_value = row)
```

Collect Stock/ETF Ticker to be displayed on future visualizations to remove confusion

```
file_name = substring(file, 6, 9)
```

First set

```
selected_stock_information_tsibble <- as_tsibble(selected_stock_information, index = row_value)
```

Set up for Monthly; second set. This could be useful in the future.

```
selected_stock_information_monthly <- yearmonth(selected_stock_information$`Date`, format = "%Y-%m-%d")
selected_stock_information$`Date` <- selected_stock_information_monthly
selected_stock_information_monthly_tsibble <- as_tsibble(selected_stock_information, index = row_value)
selected_stock_information_monthly_aggregated_tsibble <- selected_stock_information_monthly_tsibble %>% aggregate()
```

Decomposition


```
dcmp <- selected_stock_information_tsibble |>
  model(stl = STL(Close))

components(dcmp) |> autoplot() +
  labs(title = paste(file_name, "STL decomposition"),
```

```
y = "Stock Value ($)",
x = "Trading Days")
```

AAPL STL decomoposition

Close = trend + remainder




Time Series Visualization (needs another visualization)

```
autoplot(selected_stock_information_tsibble) +
  labs(y = "Stock Value ($)",
       title = paste(file_name, "Daily Closing Price over Time"),
       x = "Trading Days")
```

Plot variable not specified, automatically selected `Close`

AAPL Daily Closing Price over Time



Checking if the data is stationary. If $kpss_stat > kpss_pvalue$, then we reject null hypothesis and claim the data is non-stationary. Differencing will be required.

```
selected_stock_information_tsibble |>
  features(Close, unitroot_kpss)
```

```
## # A tibble: 1 x 2
##   kpss_stat kpss_pvalue
##   <dbl>      <dbl>
## 1      2.13         0.01
```

Output is the number of differences make the data stationary

```
selected_stock_information_tsibble |>
  features(Close, unitroot_ndiffs)
```

```
## # A tibble: 1 x 1
##   ndiffs
##   <int>
## 1     1
```

The difference

```
selected_stock_information_tsibble |>
  mutate(diff_close = difference(Close)) |>
```

```
features(diff_close, unitroot_kpss)
```

```
## # A tibble: 1 x 2
##   kpss_stat kpss_pvalue
##   <dbl>      <dbl>
## 1     0.150        0.1
```

Using Box-Cox Transformation does not change the shape of the Closing Price with a lambda value of 1.

```
lambda <- selected_stock_information_tsibble |>
  features(Close, features = guerrero) |>
  pull(lambda_guerrero)
selected_stock_information_tsibble |>
  autoplot(box_cox(Close, lambda)) +
  labs(y = "",
       title = latex2exp::TeX(paste(file_name,
                                     "Transformed Closing Price with  $\lambda = ",
                                     round(lambda, 2))))$ 
```

AAPL Transformed Closing Price with $\lambda = 1$



Description of the Time Series:

INPUT SOMETHING

TS Models (transformations)

Input SOMETHING like differencing?


Predictions

ARIMA approach

```
selected_stock_information_tsibble |>  
  gg_tsdisplay(difference(Close), plot_type='partial')
```

```
## Warning: Removed 1 row containing missing values (`geom_line()`).
```

```
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```



Generate a few ARIMA orders and then compare results.

```
selected_stock_fit <- selected_stock_information_tsibble |>  
  model(arima010 = ARIMA(Close ~ pdq(0,1,0)),  
        arima110 = ARIMA(Close ~ pdq(1,1,0)),  
        arima011 = ARIMA(Close ~ pdq(0,1,1)),  
        arima111 = ARIMA(Close ~ pdq(1,1,1)),  
        stepwise = ARIMA(Close),
```

```

search = ARIMA(Close, stepwise=FALSE))

print(selected_stock_fit$search)

## <lst_md1[1]>
## [1] <ARIMA(4,1,1)>

glance(selected_stock_fit) |> arrange(AICc) |> select(.model:BIC)

## # A tibble: 6 x 6
##   .model  sigma2 log_lik  AIC  AICc  BIC
##   <chr>    <dbl>   <dbl> <dbl> <dbl> <dbl>
## 1 search     10.2   -790. 1592. 1592. 1614.
## 2 arima010   10.4   -796. 1593. 1593. 1597.
## 3 stepwise   10.4   -796. 1593. 1593. 1597.
## 4 arima011   10.5   -796. 1595. 1595. 1603.
## 5 arima110   10.5   -796. 1595. 1595. 1603.
## 6 arima111   10.5   -795. 1597. 1597. 1608.

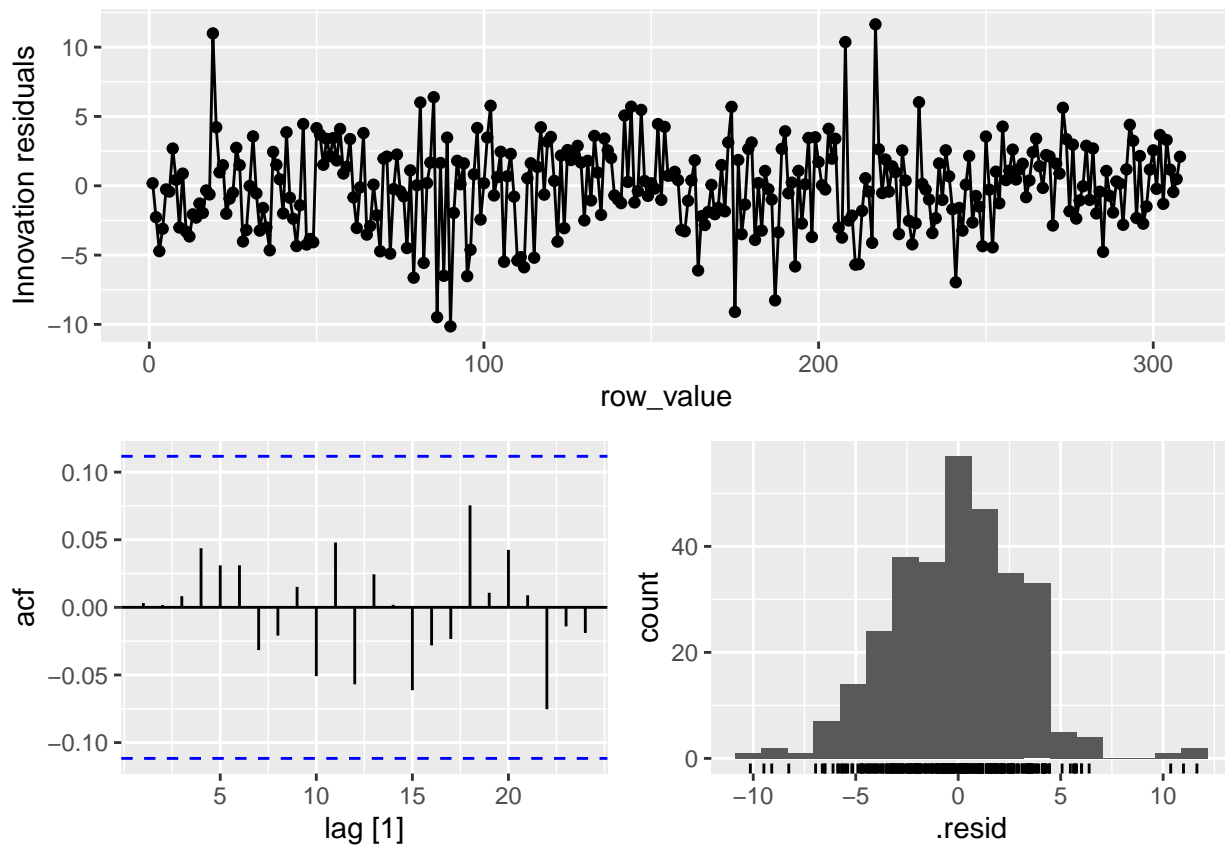
```

Show the residuals using the 'search' ARIMA

```

selected_stock_fit |>
  select(search) |>
  gg_tsresiduals()

```



Portmanteau test shows a large P-value, suggesting the residuals are similar to white noise


```
augment(selected_stock_fit) |>
  filter(.model=='search') |>
  features(.innov, ljung_box, lag = 10, dof = 3)
```

```
## # A tibble: 1 x 3
##   .model lb_stat lb_pvalue
##   <chr>    <dbl>    <dbl>
## 1 search    2.59    0.920
```

First Prediction using ARIMA

```
selected_stock_fit |>
  forecast(h=10) |>
  filter(.model=='search') |>
  autoplot(selected_stock_information_tsibble) +
  labs(y = "Stock Value ($)",
       title = paste(file_name, "Daily Closing Price over Time"),
       x = "Trading Days")
```

AAPL Daily Closing Price over Time



Second approach, first creating a Train set and a Test set

```
length_df = nrow(selected_stock_information_tsibble)
Train_number = round(length_df * 0.98, 0)
```

```
Train <- selected_stock_information_tsibble[0:Train_number,]
Test <- selected_stock_information_tsibble[Train_number:length_df,]
```

Period is set to twelve for the months, and order = 1 since data is non-seasonal

```
selected_stock_information_tsibble_fit <- Train |>
  model(
    arima = ARIMA(Close),
    ets = ETS(Close),
    prophet = prophet(Close ~ season(period = 12, order = 1,
                                     type = "additive"))
  )
```

```
## Warning: 1 error encountered for prophet
## [1] 'origin' must be supplied
```


Prophet model is not good since there is no seasonality

```
selected_stock_information_tsibble_fit_fc <- selected_stock_information_tsibble_fit |> forecast(h = 30)
selected_stock_information_tsibble_fit_fc |> autoplot() + autolayer(selected_stock_information_tsibble)
  labs(y = "Stock Value ($)",
       title = paste(file_name, "Daily Closing Price over Time"),
       x = "Trading Days")
```

```
## Plot variable not specified, automatically selected `.vars = Close`
## Warning in max(ids, na.rm = TRUE): no non-missing arguments to max; returning
## -Inf

## Warning in max(ids, na.rm = TRUE): no non-missing arguments to max; returning
## -Inf

## Warning: Removed 30 rows containing missing values (`()`).
```

MSFT

Set up File name

```
file = "Data/MSFT_2023-03-26.csv"
```

Read the .CSV and select Date and Closing Price

```
Stock_information <- read.csv(file)
selected_stock_information <- Stock_information %>% select(c('Date', 'Close')) %>% mutate(row_value = row_number())
```

Collect Stock/ETF Ticker to be displayed on future visualizations to remove confusion

```
file_name = substring(file, 6, 9)
```

First set

```
selected_stock_information_tsibble <- as_tsibble(selected_stock_information, index = row_value)
```

Set up for Monthly; second set. This could be useful in the future.

```
selected_stock_information_monthly <- yearmonth(selected_stock_information$`Date`, format = "%Y-%m-%d")
selected_stock_information$`Date` <- selected_stock_information_monthly
selected_stock_information_monthly_tsibble <- as_tsibble(selected_stock_information, index = row_value)

selected_stock_information_monthly_aggregated_tsibble <- selected_stock_information_monthly_tsibble %>% aggregate(
```

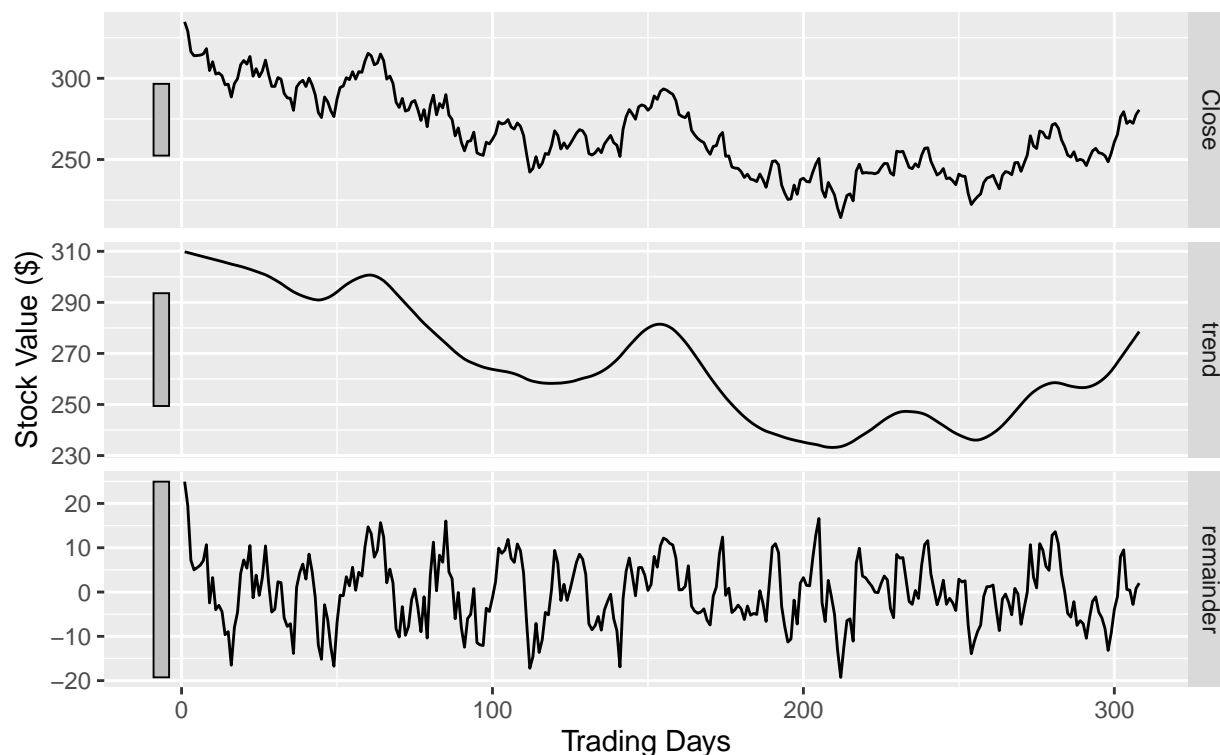
Decomposition

```
dcmp <- selected_stock_information_tsibble |>
  model(stl = STL(Close))

components(dcmp) |> autoplot() +
  labs(title = paste(file_name, "STL decomposition"),
       y = "Stock Value ($)",
       x = "Trading Days")
```

MSFT STL decomposition

Close = trend + remainder




Time Series Visualization (needs another visualization)

```
autoplot(selected_stock_information_tsibble) +
  labs(y = "Stock Value ($)",
       title = paste(file_name, "Daily Closing Price over Time"),
       x = "Trading Days")
```

Plot variable not specified, automatically selected `Close`

MSFT Daily Closing Price over Time



Checking if the data is stationary. If $\text{kpss_stat} > \text{kpss_pvalue}$, then we reject null hypothesis and claim the data is non-stationary. Differencing will be required.

```
selected_stock_information_tsibble |>
  features(Close, unitroot_kpss)
```

```
## # A tibble: 1 x 2
##   kpss_stat kpss_pvalue
##   <dbl>     <dbl>
## 1      3.37         0.01
```

Output is the number of differences make the data stationary

```
selected_stock_information_tsibble |>
  features(Close, unitroot_ndiffs)
```

```
## # A tibble: 1 x 1
##   ndiffs
##   <int>
## 1      1
```

The difference

```
selected_stock_information_tsibble |>
  mutate(diff_close = difference(Close)) |>
```

```
features(diff_close, unitroot_kpss)
```

```
## # A tibble: 1 x 2
##   kpss_stat kpss_pvalue
##   <dbl>      <dbl>
## 1      0.238        0.1
```

Using Box-Cox Transformation does not change the shape of the Closing Price with a lambda value of 1.

```
lambda <- selected_stock_information_tsibble |>
  features(Close, features = guerrero) |>
  pull(lambda_guerrero)
selected_stock_information_tsibble |>
  autoplot(box_cox(Close, lambda)) +
  labs(y = "",
       title = latex2exp::TeX(paste(file_name,
                                     "Transformed Closing Price with  $\lambda = ",
                                     round(lambda, 2))))$ 
```

MSFT Transformed Closing Price with $\lambda = 1.36$



Description of the Time Series:

INPUT SOMETHING

TS Models (transformations)

Input SOMETHING like differencing?



Predictions

ARIMA approach

```
selected_stock_information_tsibble |>  
  gg_tsdisplay(difference(Close), plot_type='partial')
```

```
## Warning: Removed 1 row containing missing values (`geom_line()`).
```

```
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```



Generate a few ARIMA orders and then compare results.

```
selected_stock_fit <- selected_stock_information_tsibble |>  
  model(arima010 = ARIMA(Close ~ pdq(0,1,0)),  
        arima110 = ARIMA(Close ~ pdq(1,1,0)),  
        arima011 = ARIMA(Close ~ pdq(0,1,1)),  
        arima111 = ARIMA(Close ~ pdq(1,1,1)),  
        stepwise = ARIMA(Close),
```

```

search = ARIMA(Close, stepwise=FALSE))

print(selected_stock_fit$search)

## <lst_mdl[1]>
## [1] <ARIMA(1,1,1)>

glance(selected_stock_fit) |> arrange(AICc) |> select(.model:BIC)

## # A tibble: 6 x 6
##   .model  sigma2 log_lik  AIC  AICc  BIC
##   <chr>    <dbl>   <dbl> <dbl> <dbl> <dbl>
## 1 arima010  33.1   -973.  1948.  1948.  1952.
## 2 arima111  33.1   -972.  1950.  1950.  1961.
## 3 stepwise  33.1   -972.  1950.  1950.  1961.
## 4 search    33.1   -972.  1950.  1950.  1961.
## 5 arima011  33.2   -973.  1950.  1950.  1957.
## 6 arima110  33.2   -973.  1950.  1950.  1957.

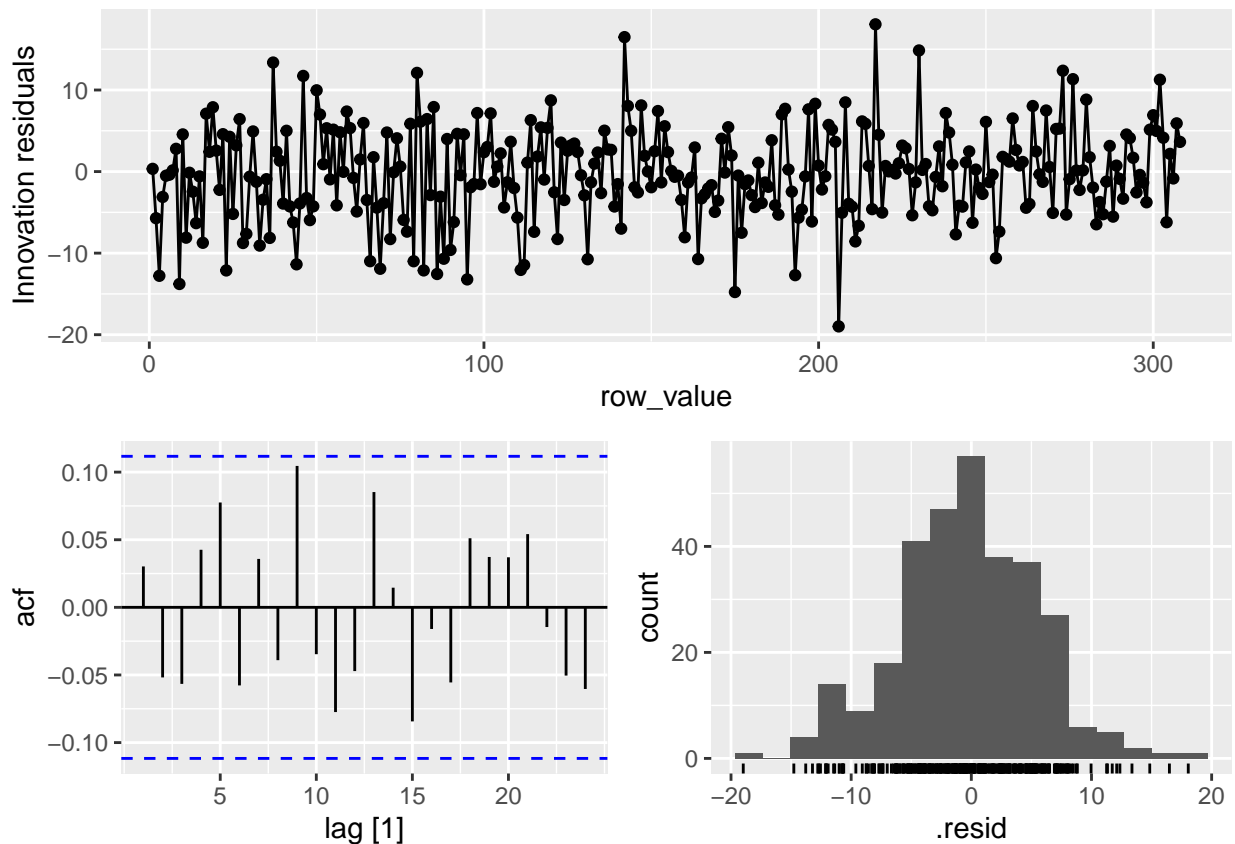
```

Show the residuals using the 'search' ARIMA

```

selected_stock_fit |>
  select(search) |>
  gg_tsresiduals()

```



Portmanteau test shows a large P-value, suggesting the residuals are similar to white noise

```
augment(selected_stock_fit) |>
  filter(.model=='search') |>
  features(.innov, ljung_box, lag = 10, dof = 3)
```

```
## # A tibble: 1 x 3
##   .model lb_stat lb_pvalue
##   <chr>    <dbl>    <dbl>
## 1 search    10.4      0.166
```

First Prediction using ARIMA

```
selected_stock_fit |>
  forecast(h=10) |>
  filter(.model=='search') |>
  autoplot(selected_stock_information_tsibble) +
  labs(y = "Stock Value ($)",
       title = paste(file_name, "Daily Closing Price over Time"),
       x = "Trading Days")
```

MSFT Daily Closing Price over Time



Second approach, first creating a Train set and a Test set

```
length_df = nrow(selected_stock_information_tsibble)
Train_number = round(length_df * 0.98, 0)
```

```
Train <- selected_stock_information_tsibble[0:Train_number,]
Test <- selected_stock_information_tsibble[Train_number:length_df,]
```

Period is set to twelve for the months, and order = 1 since data is non-seasonal

```
selected_stock_information_tsibble_fit <- Train |>
  model(
    arima = ARIMA(Close),
    ets = ETS(Close),
    prophet = prophet(Close ~ season(period = 12, order = 1,
                                     type = "additive"))
  )
```

```
## Warning: 1 error encountered for prophet
## [1] 'origin' must be supplied
```

Prophet model is not good since there is no seasonality


```
selected_stock_information_tsibble_fit_fc <- selected_stock_information_tsibble_fit |> forecast(h = 30)
selected_stock_information_tsibble_fit_fc |> autoplot() + autolayer(selected_stock_information_tsibble)
  labs(y = "Stock Value ($)",
       title = paste(file_name, "Daily Closing Price over Time"),
       x = "Trading Days")
```

```
## Plot variable not specified, automatically selected `.vars = Close`
## Warning in max(ids, na.rm = TRUE): no non-missing arguments to max; returning
## -Inf

## Warning in max(ids, na.rm = TRUE): no non-missing arguments to max; returning
## -Inf

## Warning: Removed 30 rows containing missing values (`()`).
```


MSFT Daily Closing Price over Time



TSLA

Set up File name

```
file = "Data/TSLA_2023-03-26.csv"
```

Read the .CSV and select Date and Closing Price

```
Stock_information <- read.csv(file)
selected_stock_information <- Stock_information %>% select(c('Date','Close')) %>% mutate(row_value = row_number())
```

Collect Stock/ETF Ticker to be displayed on future visualizations to remove confusion

```
file_name = substring(file,6,9)
```

First set

```
selected_stock_information_tsibble <- as_tsibble(selected_stock_information, index = row_value)
```

Set up for Monthly; second set. This could be useful in the future.


```
selected_stock_information_monthly <- yearmonth(selected_stock_information$`Date`, format = "%Y-%m-%d")
selected_stock_information$`Date` <- selected_stock_information_monthly
selected_stock_information_monthly_tsibble <- as_tsibble(selected_stock_information, index = row_value)

selected_stock_information_monthly_aggregated_tsibble <- selected_stock_information_monthly_tsibble %>% aggregate(
```

Decomposition

```
dcmp <- selected_stock_information_tsibble |>
  model(stl = STL(Close))

components(dcmp) |> autoplot() +
  labs(title = paste(file_name, "STL decomposition"),
       y = "Stock Value ($)",
       x = "Trading Days")
```



Time Series Visualization (needs another visualization)

```
autoplot(selected_stock_information_tsibble) +
  labs(y = "Stock Value ($)",
       title = paste(file_name, "Daily Closing Price over Time"),
       x = "Trading Days")
```

Plot variable not specified, automatically selected `Close`

TSLA Daily Closing Price over Time



Checking if the data is stationary. If $kpss_stat > kpss_pvalue$, then we reject null hypothesis and claim the data is non-stationary. Differencing will be required.

```
selected_stock_information_tsibble |>
  features(Close, unitroot_kpss)
```

```
## # A tibble: 1 x 2
##   kpss_stat kpss_pvalue
##   <dbl>     <dbl>
## 1      3.58         0.01
```

Output is the number of differences make the data stationary

```
selected_stock_information_tsibble |>
  features(Close, unitroot_ndiffs)
```

```
## # A tibble: 1 x 1
##   ndiffs
##   <int>
## 1     1
```

The difference

```
selected_stock_information_tsibble |>
  mutate(diff_close = difference(Close)) |>
```

```
features(diff_close, unitroot_kpss)
```

```
## # A tibble: 1 x 2
##   kpss_stat kpss_pvalue
##   <dbl>      <dbl>
## 1     0.127        0.1
```

Using Box-Cox Transformation does not change the shape of the Closing Price with a lambda value of 1.

```
lambda <- selected_stock_information_tsibble |>
  features(Close, features = guerrero) |>
  pull(lambda_guerrero)
selected_stock_information_tsibble |>
  autoplot(box_cox(Close, lambda)) +
  labs(y = "",
       title = latex2exp::TeX(paste(file_name,
                                     "Transformed Closing Price with  $\lambda = ",
                                     round(lambda, 2))))$ 
```

TSLA Transformed Closing Price with $\lambda = 0.52$



Description of the Time Series:

INPUT SOMETHING

TS Models (transformations)

Input SOMETHING like differencing?

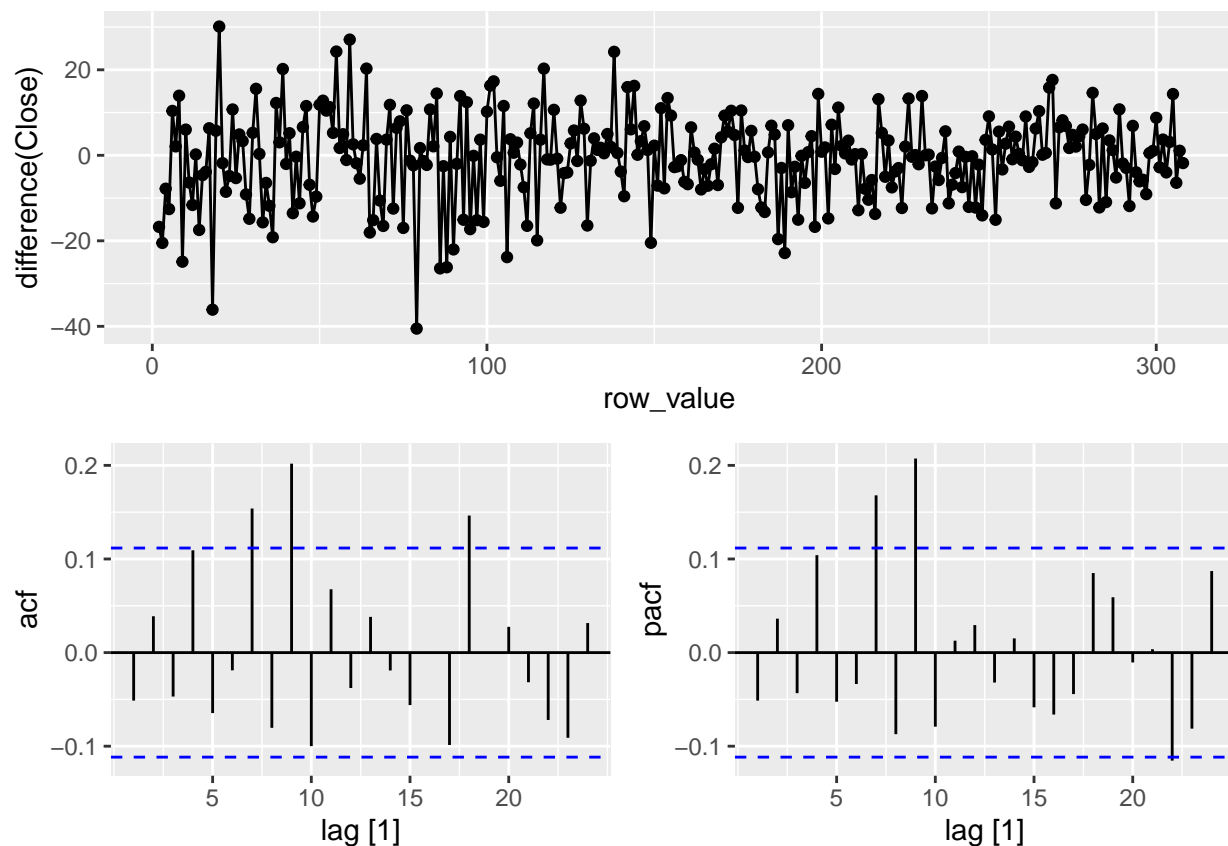
Predictions

ARIMA approach

```
selected_stock_information_tsibble |>  
  gg_tsdisplay(difference(Close), plot_type='partial')
```

```
## Warning: Removed 1 row containing missing values (`geom_line()`).
```

```
## Warning: Removed 1 rows containing missing values (`geom_point()`).
```



Generate a few ARIMA orders and then compare results.

```
selected_stock_fit <- selected_stock_information_tsibble |>  
  model(arima010 = ARIMA(Close ~ pdq(0,1,0)),  
        arima110 = ARIMA(Close ~ pdq(1,1,0)),  
        arima011 = ARIMA(Close ~ pdq(0,1,1)),  
        arima111 = ARIMA(Close ~ pdq(1,1,1)),  
        stepwise = ARIMA(Close),
```

```

search = ARIMA(Close, stepwise=FALSE))

print(selected_stock_fit$search)

## <lst_md1[1]>
## [1] <ARIMA(0,1,5)>

glance(selected_stock_fit) |> arrange(AICc) |> select(.model:BIC)

## # A tibble: 6 x 6
##   .model  sigma2 log_lik  AIC  AICc  BIC
##   <chr>    <dbl>   <dbl> <dbl> <dbl> <dbl>
## 1 search      102.  -1143.  2297.  2298.  2320.
## 2 arima010    105.  -1150.  2301.  2301.  2305.
## 3 arima110    105.  -1149.  2302.  2302.  2310.
## 4 arima011    105.  -1149.  2302.  2302.  2310.
## 5 stepwise    105.  -1149.  2304.  2304.  2315.
## 6 arima111    105.  -1149.  2305.  2305.  2316.

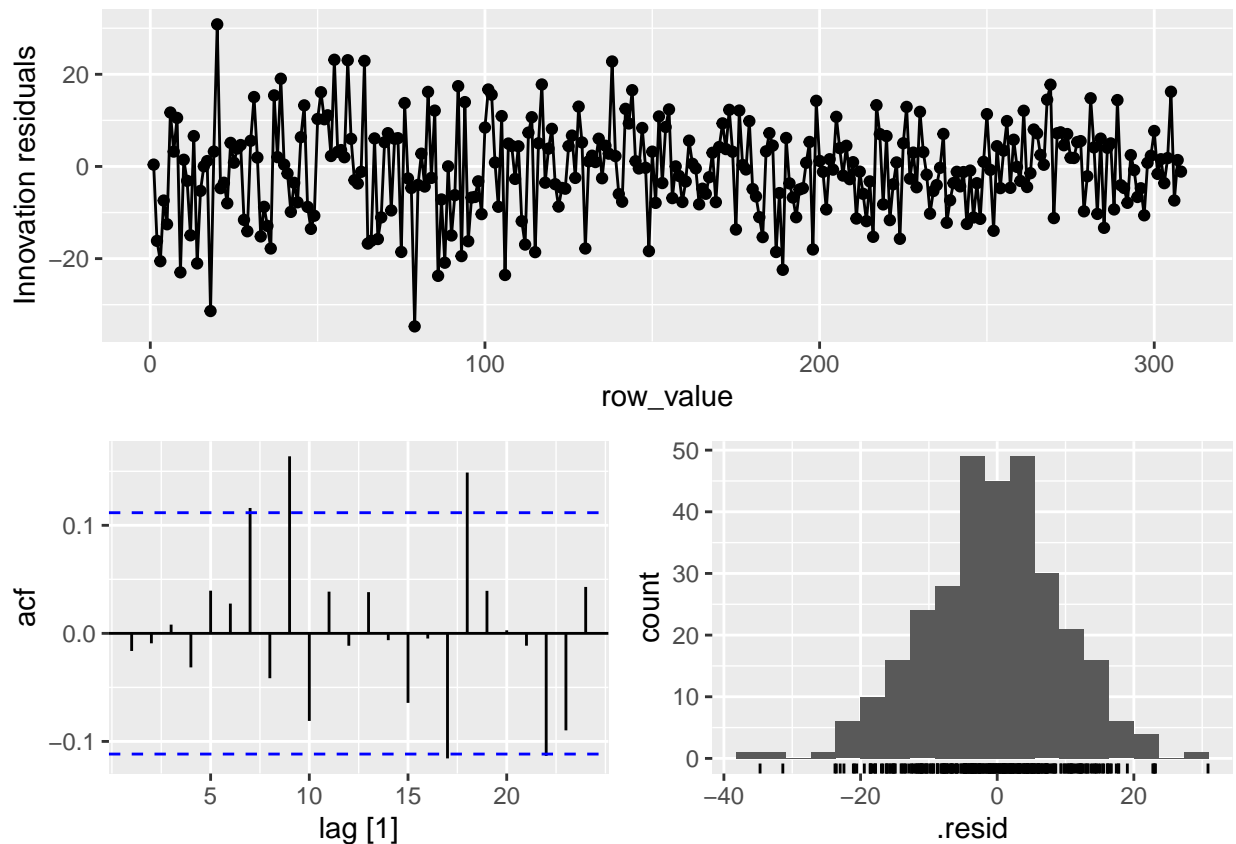
```

Show the residuals using the 'search' ARIMA

```

selected_stock_fit |>
  select(search) |>
  gg_tsresiduals()

```




Portmanteau test shows a large P-value, suggesting the residuals are similar to white noise

```
augment(selected_stock_fit) |>
  filter(.model=='search') |>
  features(.innov, ljung_box, lag = 10, dof = 3)
```

```
## # A tibble: 1 x 3
##   .model lb_stat lb_pvalue
##   <chr>   <dbl>   <dbl>
## 1 search    16.7    0.0197
```

First Prediction using ARIMA

```
selected_stock_fit |>
  forecast(h=10) |>
  filter(.model=='search') |>
  autoplot(selected_stock_information_tsibble) +
  labs(y = "Stock Value ($)",
       title = paste(file_name, "Daily Closing Price over Time"),
       x = "Trading Days")
```



Second approach, first creating a Train set and a Test set

```
length_df = nrow(selected_stock_information_tsibble)
Train_number = round(length_df * 0.98, 0)
```

```
Train <- selected_stock_information_tsibble[0:Train_number,]
Test <- selected_stock_information_tsibble[Train_number:length_df,]
```

Period is set to twelve for the months, and order = 1 since data is non-seasonal

```
selected_stock_information_tsibble_fit <- Train |>
  model(
    arima = ARIMA(Close),
    ets = ETS(Close),
    prophet = prophet(Close ~ season(period = 12, order = 1,
                                     type = "additive"))
  )
```

```
## Warning: 1 error encountered for prophet
## [1] 'origin' must be supplied
```


Prophet model is not good since there is no seasonality

```
selected_stock_information_tsibble_fit_fc <- selected_stock_information_tsibble_fit |> forecast(h = 30)
selected_stock_information_tsibble_fit_fc |> autoplot() + autolayer(selected_stock_information_tsibble)
  labs(y = "Stock Value ($)",
       title = paste(file_name, "Daily Closing Price over Time"),
       x = "Trading Days")
```

```
## Plot variable not specified, automatically selected `.vars = Close`
## Warning in max(ids, na.rm = TRUE): no non-missing arguments to max; returning
## -Inf

## Warning in max(ids, na.rm = TRUE): no non-missing arguments to max; returning
## -Inf

## Warning: Removed 30 rows containing missing values (`()`).
```

Team contributions (pending...)

Brian: Completed rough draft of the Part 2 submission.

Mohan: Generated Github repo and place Datasets. Picked option 1.

Brendan: Picked option 1.